



# 獨孤派作業系統 期中，上機考

中正大學 作業系統實驗室  
指導教授：羅習五



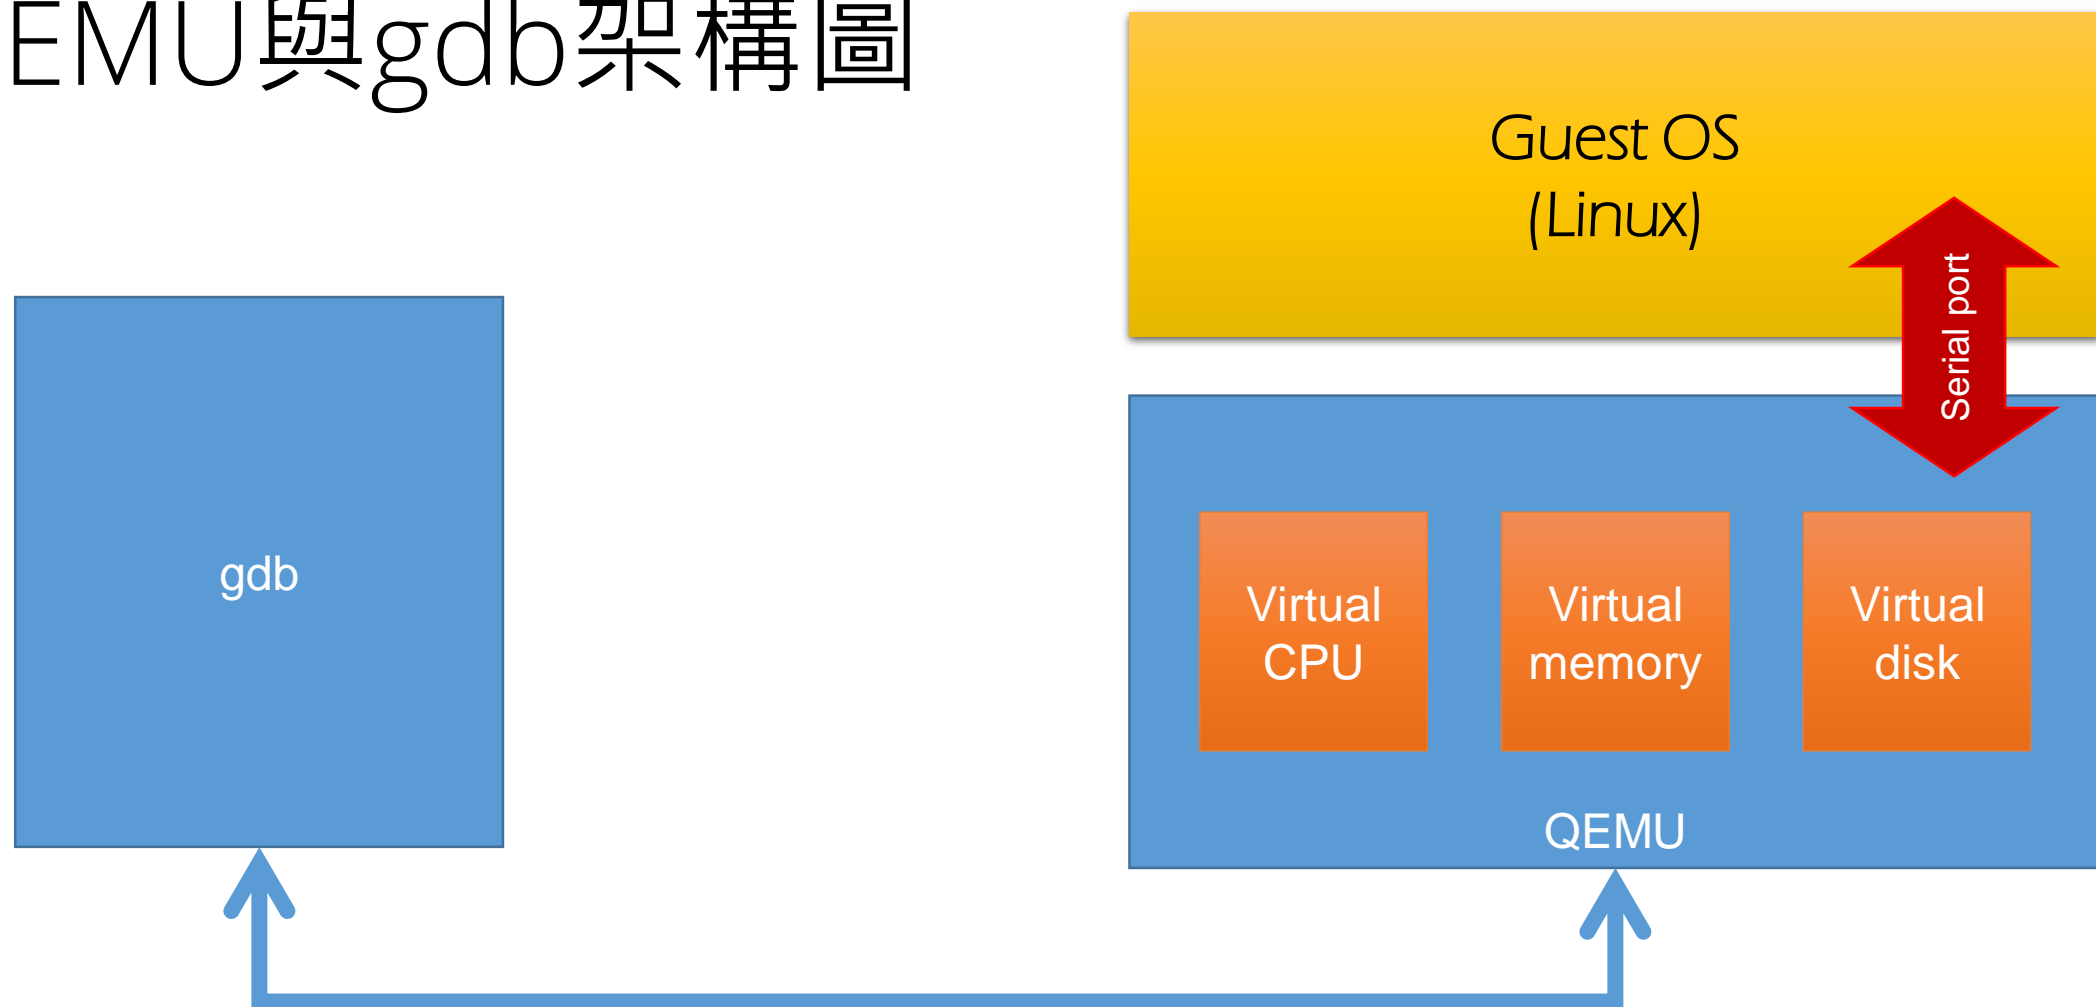
# 下載已經做好的qemu相關檔案

- 自dropbox下載home\_shiwulo-qemu+eclipse+oracleJava.tar.xz
- 解壓縮
  - `sudo tar Jxvf home_shiwulo-qemu+eclipse+oracleJava.tar.xz`
  - 注意，這個地方要用sudo，因為壓縮檔中有特殊檔案，需要用mknode產生
  - 注意，解壓縮以後，這些檔案的owner和group都是「1000」，如果你的id不是1000請使用chown及chgrp修改

# 目錄內容介紹

- 所有相關檔案都放在/home/yourname/qemu-linux-4.0底下
- 『dbg-qemu.sh』使用qemu執行Linux，qemu會等待debugger連入
- 『vmlinux』是帶有debug資訊的Linux的核心
- 『bzImage、initramfs\_data.cpio.gz』啟動用的Linux kernel及initramfs
- 『system.map』使用『nm』自『vmlinux』匯出所有的symbol的可讀檔案

# QEMU與gdb架構圖



很特別，這個地方是走TCP/IP，因此可以（如果需要的話）  
可以遠端除錯

# 開始本次上機考



# 啟動debuggee

- `$ ./ dbg-qemu.sh`
- 啟動以後，QEMU會停住，等待gdb的連線



# 啟動debugger

- 另外開一個terminal，輸入
- `$ gdb ./vmlinux`
- gdb會讀取vmlinux的symbol table，進入gdb後...
- `(gdb) target remote localhost:666`



# 第一部分，評分方式

1. ( 15pt ) 請問你是用懶人包，或者自行設定trace kernel的環境？
  - a) ( 15pt ) 完全自行設定
  - b) ( 10pt ) 使用/home/shiwulo的懶人包
  - c) ( 5pt ) 使用virtual machine懶人包
2. ( 10pt ) 設定breakpoint於start\_kernel，並印出start\_kernel的call stack，及call stack中各個函數的位址
3. ( 10pt ) 找出中斷向量表的開始位址
4. ( 10pt ) 找出system call的進入點
5. ( 10pt ) 印出CPU初始化的時候，「除以零」這個中斷的interrupt service routine的位址



## 2. start\_kernel

這個號碼就是第  
#題的答案

#0 start\_kernel () at init/main.c:490

#1 0xffffffff822db72d in x86\_64\_start\_reservations (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200

#2 0xffffffff822db6f0 in x86\_64\_start\_kernel (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189

#3 0x0000000000000000 in ?? ()

### 3. IDT進入點（即中斷向量表開始位址）

```
(gdb) p idt_table
```

```
$3 = 0xffffffff825fb000 <idt_table>
```

```
(gdb) bt
```

```
#0 _set_gate (gate=<optimized out>, type=<optimized out>, addr=0xffffffff81d0f850 <divide_error>, dpl=<optimized out>, ist=<optimized out>, seg=<optimized out>) at ./arch/x86/include/asm/desc.h:371
```

```
#1 0xffffffff822df55c in trap_init () at arch/x86/kernel/traps.c:956
```

```
#2 0xffffffff822dc07b in start_kernel () at init/main.c:550
```

```
#3 0xffffffff822db72d in x86_64_start_reservations (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200
```

```
#4 0xffffffff822db6f0 in x86_64_start_kernel (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189
```

```
#5 0x0000000000000000 in ?? ()
```

```
(gdb) p/x *(struct gate_struct64*)idt_table
```

```
$6 = {offset_low = 0xb000, segment = 0x10, ist = 0x0, zero0 = 0x0, type = 0xe, dpl = 0x0, p = 0x1, offset_middle = 0x822d, offset_high = 0xffffffff, zero1 = 0x0}
```

# TSS ( cpu0 )

(gdb) p t

\$4 = (struct tss\_struct \*) 0xffff88000fa10ec0

(gdb) p \*t

\$5 = {x86\_tss = {reserved1 = 0, sp0 = 18446744071597735936, sp1 = 0, sp2 = 0, reserved2 = 0, ist = {0, 0, 0, 0, 0, 0, 0, 0}, reserved3 = 0, reserved4 = 0, reserved5 = 0, io\_bitmap\_base = 0, io\_bitmap = {0 <repeats 1025 times>}, stack = {0 <repeats 64 times>}}

(gdb) bt

#0 cpu\_init () at arch/x86/kernel/cpu

#1 0xffffffff822dfb80 in trap\_init () at arch/x86/kernel/traps.c:1006

#2 0xffffffff822dc07b in start\_kernel () at init/main.c:550

#3 0xffffffff822db72d in x86\_64\_start\_reservations (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200

#4 0xffffffff822db6f0 in x86\_64\_start\_kernel (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189

#5 0x0000000000000000 in ?? ()

這個值很重要，每個CPU  
有自己的TSS

0xFFFFFFFF82204000

每個task有自己的stack即sp0

# idt descriptor

```
(gdb) p/x idt_descr
```

```
$13 = {size = 0xfff, address = 0xffffffff57b000}
```

```
(gdb) bt
```

```
#0 load_current_idt () at ./arch/x86/include/asm/desc.h:513
```

```
#1 0xffffffff810230d8 in cpu_init () at arch/x86/kernel/cpu/common.c:1333
```

```
#2 0xffffffff822dfb80 in trap_init () at arch/x86/kernel/traps.c:1006
```

```
#3 0xffffffff822dc07b in start_kernel () at init/main.c:550
```

```
#4 0xffffffff822db72d in x86_64_start_reservations (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200
```

```
#5 0xffffffff822db6f0 in x86_64_start_kernel (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189
```

```
#6 0x0000000000000000 in ?? ()
```

## 4. sytem call

(gdb) p system\_call

\$15 = {<text variable, no debug info>} 0xffffffff81d0e100 <system\_call>

(gdb) bt

#0 syscall\_init () at arch/x86/kernel/cpu/common.c:1174

#1 0xffffffff810230f9 in cpu\_init () at arch/x86/kernel/cpu/common.c:1336

#2 0xffffffff822dfb80 in trap\_init () at arch/x86/kernel/traps.c:1006

#3 0xffffffff822dc07b in start\_kernel () at init/main.c:550

#4 0xffffffff822db72d in x86\_64\_start\_reservations (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200

#5 0xffffffff822db6f0 in x86\_64\_start\_kernel (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189

#6 0x0000000000000000 in ?? ()

# interrupt stack

(gdb) p/x \*t

\$21 = {x86\_tss = {reserved1 = 0x0, sp0 = 0xffffffff82204000, sp1 = 0x0, sp2 = 0x0, reserved2 = 0x0, ist = {0xffff88000fa06000, 0xffff88000fa07000, 0xffff88000fa09000, 0xffff88000fa0a000, 0x0, 0x0, 0x0}, reserved3 = 0x0, reserved4 = 0x0, reserved5 = 0x0, io\_bitmap\_base = 0x0}, io\_bitmap = {0x0 <repeats 1025 times>}, stack = {0x0 <repeats 64 times>}}

(gdb) bt

#0 cpu\_init () at arch/x86/kernel/cpu/common.c:1360

#1 0xffffffff822dfb80 in trap\_init () at arch/x86/kernel/traps.c:1006

#2 0xffffffff822dc07b in start\_kernel () at init/main.c:550

#3 0xffffffff822db72d in x86\_64\_start\_reservations (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200

#4 0xffffffff822db6f0 in x86\_64\_start\_kernel (real\_mode\_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189

#5 0x0000000000000000 in ?? ()

# interrupt stack

```
(gdb) p/x *(struct tss_struct*) 0xffff88000fa10ec0
```

```
$13 = {x86_tss = {reserved1 = 0x0, sp0 = 0xffffffff82204000, sp1 = 0x0, sp2 = 0x0, reserved2 = 0x0, ist =  
{0xffff88000fa06000, 0xffff88000fa07000, 0xffff88000fa09000, 0xffff88000fa0a000, 0x0, 0x0, 0x0}, reserved3 =  
0x0, reserved4 = 0x0, reserved5 = 0x0, io_bitmap_base = 0x80}, io_bitmap = {0xffffffffffffffff <repeats 1025 times>},  
stack = {0x0 <repeats 64 times>}}
```

```
(gdb) bt
```

```
#0 0xffffffff81004416 in __switch_to (prev_p=0xffffffff8220e180 <init_task>, next_p=0xffff88000ecc0910) at  
arch/x86/kernel/process_64.c:280
```

```
#1 0xffffffff810c8e52 in context_switch (rq=0x0 <irq_stack_union>, prev=0x0 <irq_stack_union>, next=0x0  
<irq_stack_union>) at kernel/sched/core.c:2317
```

```
(gdb) p/x *(struct gate_struct64*) 0xffffffff825fb000
```

```
$14 = {offset_low = 0xf850, segment = 0x10, ist = 0x0, zero0 = 0x0, type = 0xe, dpl = 0x0, p = 0x1, offset_middle =  
0x81d0, offset_high = 0xffffffff, zero1 = 0x0}
```


## 5. 印出第零號中斷 即「除以零」

「除以零的ISR位址」 `ffffffff, 81d0, f850`

```
(gdb) p/x *(struct gate_struct64*) 0xffffffff825fb000
```

```
$14 = {offset_low = 0xf850, segment = 0x10, ist = 0x0, zero0 =  
0x0, type = 0xe, dpl = 0x0, p = 0x1, offset_middle = 0x81d0,  
offset_high = 0xffffffff, zero1 = 0x0}
```

資料結構的定義請見  
下一頁





# gate\_struct64

## /arch/x86/include/asm/desc\_defs.h

```
1. /* 16byte gate */
2. struct gate_struct64 {
3.     u16 offset_low;
4.     u16 segment;
5.     unsigned ist : 3, zero0 : 5, type : 5, dpl : 2, p : 1;
6.     u16 offset_middle;
7.     u32 offset_high;
8.     u32 zero1;
9. } __attribute__((packed));
```

# system call所用的kernel stack位置

## 每一個task會有自己的kernel stack

```
(gdb) p/x thread->sp0
```

```
$24 = 0xffffffff82204000
```

```
(gdb) bt
```

```
#0 0xffffffff82204000 in ?? at ./a.out:0, thread=0xffffffff8220e778 <init_task+1528>)
```

```
at ./a.out:58
```

```
#1 0xffffffff82204000 in ?? at ./a.out:76, thread=0xffffffff8220e778 <init_task+1528>)
```

```
at ./a.out:76
```

```
#2 0xffffffff810232a3 in cpu_init () at arch/x86/kernel/cpu/common.c:1374
```

```
#3 0xffffffff822dfb80 in trap_init () at arch/x86/kernel/traps.c:1006
```

```
#4 0xffffffff822dc07b in start_kernel () at init/main.c:550
```

```
#5 0xffffffff822db72d in x86_64_start_reservations (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:200
```

```
#6 0xffffffff822db6f0 in x86_64_start_kernel (real_mode_data=0x13f60 <runqueues+1440> <error: Cannot access memory at address 0x13f60>) at arch/x86/kernel/head64.c:189
```

```
#7 0x0000000000000000 in ?? ()
```

跟第12頁的值一樣

0xFFFFFFFF82204000

## 第二部分，評分方式

6. ( 30pt ) 於「執行」 divideByZero時，使用gdb對中斷處理常式 ( interrupt service routine，ISR的位址 ) 進行除錯
  - a) ( 10pt ) 印出當下正在執行的ISR的位址
  - b) ( 10pt ) 印出ISR的程式碼 ( 組合語言 )
  - c) ( 10pt ) 印出該ISR所使用的堆疊的位址，請問此時ISR所使用的stack是從哪邊載入
7. ( 15pt ) ( 延續問題5.a ) 請解釋ISR組合語言的意義
  - a) ( 10pt ) error\_entry
  - b) ( 5pt ) do\_divide\_error



# b divide\_error

(gdb) # 按下ctr-c

(gdb) b divide\_error

Breakpoint 12 at 0xffffffff81d0f850: file  
arch/x86/kernel/entry\_64.S, line 1020.

(gdb) c

# 在QEMU中

/ # ./divideByZero

## 6.a、6.b 第一次攔截「除以零」 ISR如下所示

(gdb) disassemble /m divide\_error

Dump of assembler code for function divide\_error:

```
1143  idtentry divide_error do_divide_error has_error_code=0
0xffffffff81d0f853 <+3>:  pushq $0xffffffffffffff
0xffffffff81d0f855 <+5>:  sub   $0x78,%rsp
0xffffffff81d0f859 <+9>:  callq 0xffffffff81d0fd20 <error_entry>
0xffffffff81d0f85e <+14>: mov   %rsp,%rdi
0xffffffff81d0f861 <+17>: xor   %esi,%esi
0xffffffff81d0f863 <+19>: callq 0xffffffff8100688f <do_divide_error>
0xffffffff81d0f868 <+24>: jmpq  0xffffffff81d0fdd0 <error_exit>
0xffffffff81d0f86d: nopl  (%rax)
0xffffffff81d0f870 <+0>:  data16 xchg %ax,%ax
```

End of assembler dump.

# 第一次攔截「除以零」

## 印出暫存器

```
(gdb) info registers
rax          0xa          10
rbx          0x400400    4195328
rcx          0x44b9c0    4504000
rdx          0x0          0
rsi          0x7fff252c40a8
rdi          0x1          1
rbp          0x7fff252c3f80
rsp          0xffff88000002bfd8
r8           0x1000000    16777216
r9           0x6ba8e0     7055584
r10          0x15         21
r11          0x0          0
r12          0x401900    4200704
r13          0x0          0
r14          0x6b8018    7045144
r15          0x0          0
rip          0xffffffff81d0f850
<divide_error>
eflags       0x46        [ PF ZF ]
cs           0x10        16
ss           0x0          0
ds           0x0          0
es           0x0          0
fs           0x63         99
gs           0x0          0
```

## 6.c 第二次攔截「除以零」

```
(gdb) info registers
```

rax	0xa	10	r12	0x401900	4200704
rbx	0x400400	4195328	r13	0x0	0
rcx	0x44b9c0	4504000	r14	0x6b8018	7045144
rdx	0x0	0	r15	0x0	0
rsi	0x7fff7fba26b8		rip	0xffffffff81d0f850	
rdi	0x1	1	<divide_error>		
rbp	0x7fff7fba2590		eflags	0x46	[ PF ZF ]
rsp	0xffff8800e59bfd8		cs	0x10	16
r8	0x1000000	16777216			0
r9	0x6ba8e0	7055584	es	0x0	0
r10	0x15	21	fs	0x63	99
r11	0x0	0	gs	0x0	0

rsp不一樣是因為每個task  
有自己的kernel stack



## 6.c 第二次攔截「除以零」 觀察TSS的結構

0xffff88000e59bfd8 -  
0xffff88000e59c000 =  
0x28

「除以零」和「system call」  
都屬於exception，在x86中堆  
疊會自動切換到sp0

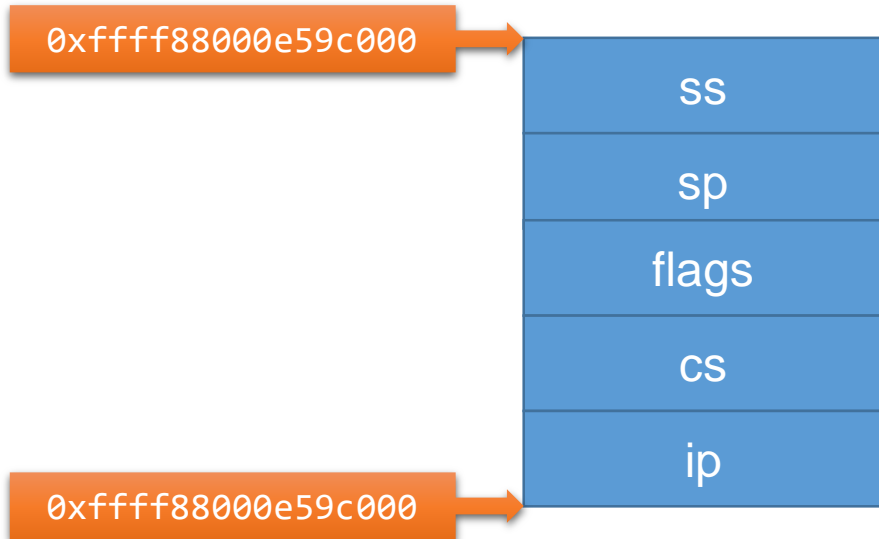
(gdb) p/x \*(struct tss\_struct \*) 0xffff88000fa10ec0

\$15 = {x86\_tss = {reserved1 = 0x0, sp0 = 0xffff88000e59c000, sp1 = 0x0, sp2 = 0x0, reserved2 = 0x0, ist = {0xffff88000fa06000, 0xffff88000fa07000, 0xffff88000fa09000, 0xffff88000fa0a000, 0x0, 0x0, 0x0}, reserved3 = 0x0, reserved4 = 0x0, reserved5 = 0x0, io\_bitmap\_base = 0x80}, io\_bitmap = {0xffffffffffffffff <repeats 1025 times>}, stack = {0x0 <repeats 64 times>}}

## 7.a 關於x86-64如何處理exception的文章

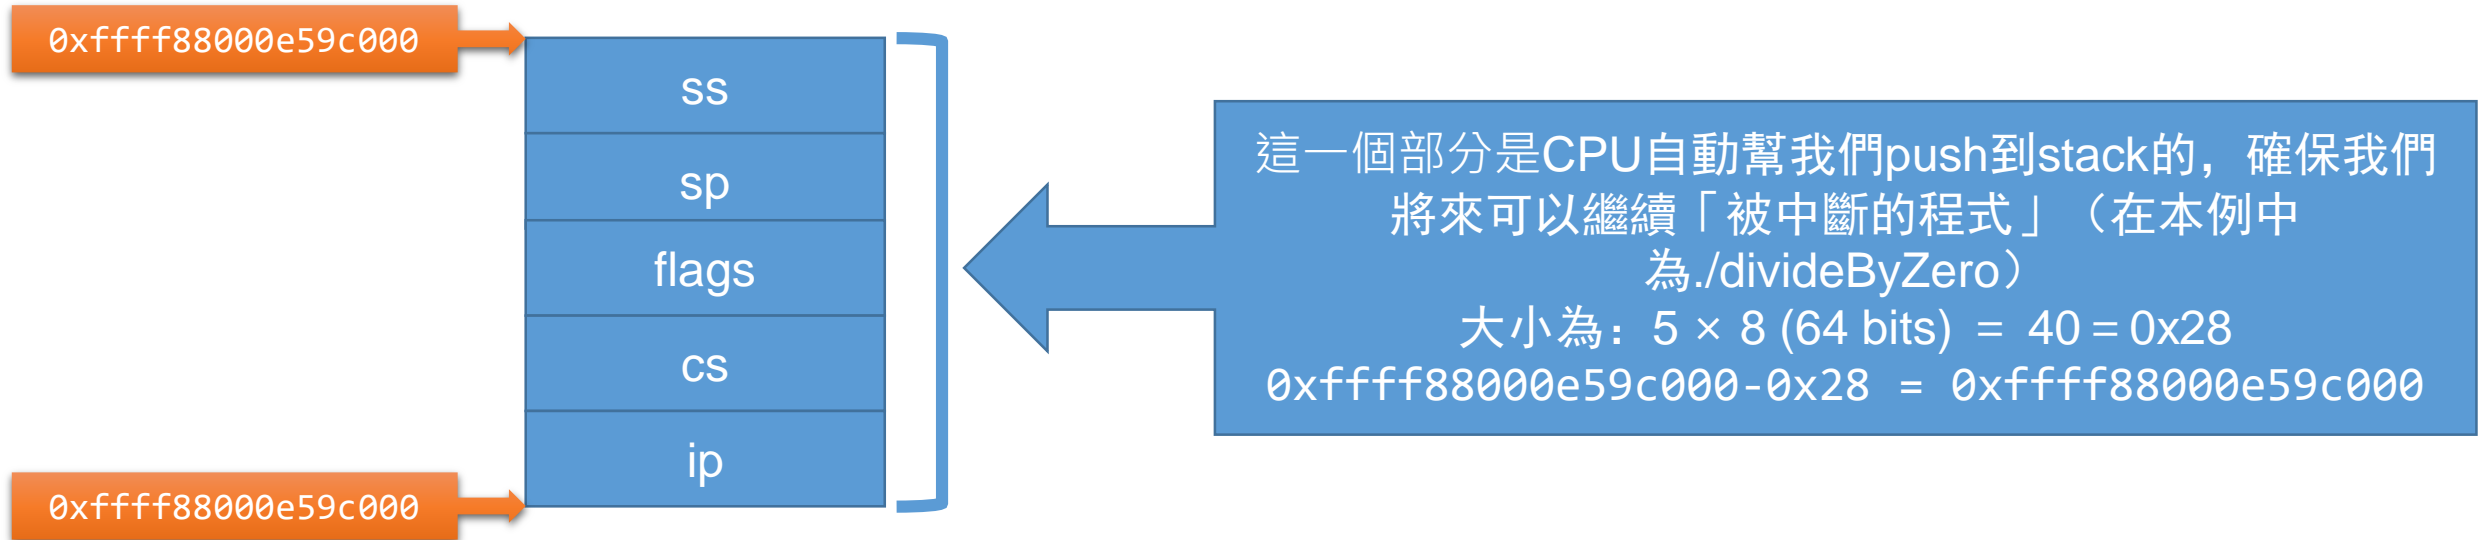
- **Aligning the stack pointer:** An interrupt can occur at any instructions, so the stack pointer can have any value, too. However, some CPU instructions (e.g. some SSE instructions) require that the stack pointer is aligned on a 16 byte boundary, therefore the CPU performs such an alignment right after the interrupt.
- **Switching stacks** (in some cases): A stack switch occurs when the CPU privilege level changes, for example when a CPU exception occurs in an user mode program. It is also possible to configure stack switches for specific interrupts using the so-called *Interrupt Stack Table* (described in the next post).
- **Pushing the old stack pointer:** The CPU pushes the values of the stack pointer (**rsp**) and the stack segment (**ss**) registers at the time when the interrupt occurred (before the alignment). This makes it possible to restore the original stack pointer when returning from an interrupt handler.
- **Pushing and updating the RFLAGS register:** The **RFLAGS** register contains various control and status bits. On interrupt entry, the CPU changes some bits and pushes the old value.
- **Pushing the instruction pointer:** Before jumping to the interrupt handler function, the CPU pushes the instruction pointer (**rip**) and the code segment (**cs**). This is comparable to the return address push of a normal function call.
- **Pushing an error code** (for some exceptions): For some specific exceptions such as page faults, the CPU pushes an error code, which describes the cause of the exception.
- **Invoking the interrupt handler:** The CPU reads the address and the segment descriptor of the interrupt handler function from the corresponding field in the IDT. It then invokes this handler by loading the values into the rip and cs registers.

## 7.a 堆疊的宣告



```
struct pt_regs {  
    unsigned long r15;  
    unsigned long r14;  
    unsigned long r13;  
    unsigned long r12;  
    unsigned long bp;  
    unsigned long bx;  
  
    /* arguments: non interrupts/non tracing syscalls only save up to  
    here*/  
    unsigned long r11;  
    unsigned long r10;  
    unsigned long r9;  
    unsigned long r8;  
    unsigned long ax;  
    unsigned long cx;  
    unsigned long dx;  
    unsigned long si;  
    unsigned long di;  
    unsigned long orig_ax;  
  
    /* end of arguments */  
    /* cpu exception frame or undefined */  
    unsigned long ip;  
    unsigned long cs;  
    unsigned long flags;  
    unsigned long sp;  
    unsigned long ss;  
  
    /* top of stack page */  
};
```

## 7.a 堆疊的宣告



如果是system call這邊會放進去AX, exception則放「錯誤代碼」

## 7.a 堆疊的宣告

0xffff88000e59c000

ss

sp

flags

cs

ip

0xffff88000e59c000

orig\_ax

↓

r15

```
pushq    $0xffffffffffffffff
sub      $0x78,%rsp
callq    0xfffffffff81d0fd20 <error_entry>
mov      %rsp,%rdi
xor      %esi,%esi
callq    0xfffffffff8100688f <do_divide_error>
jmpq     0xfffffffff81d0fdd0 <error_exit>
nopl     (%rax)
data16   xchg %ax,%ax
```

cld

```
mov      %rdi,0x78(%rsp)
mov      %rsi,0x70(%rsp)
mov      %rdx,0x68(%rsp)
mov      %rcx,0x60(%rsp)
mov      %rax,0x58(%rsp)
mov      %r8,0x50(%rsp)
mov      %r9,0x48(%rsp)
mov      %r10,0x40(%rsp)
mov      %r11,0x38(%rsp)
mov      %rbx,0x30(%rsp)
mov      %rbp,0x28(%rsp)
mov      %r12,0x20(%rsp)
mov      %r13,0x18(%rsp)
mov      %r14,0x10(%rsp)
mov      %r15,0x8(%rsp)
```

見下頁解釋

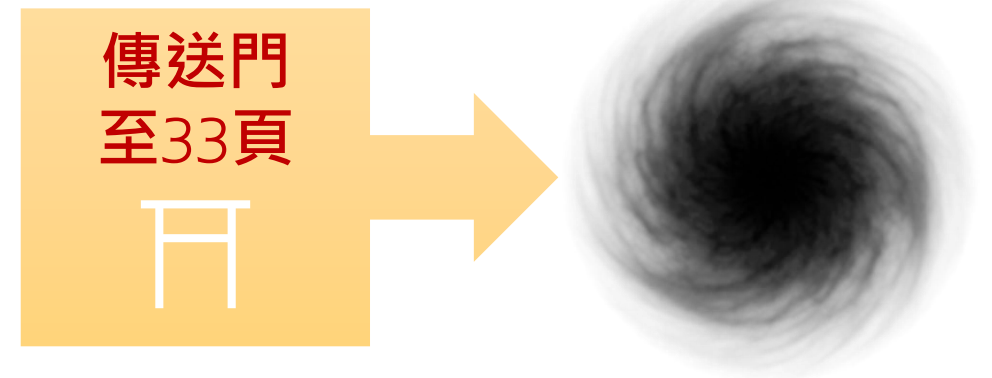
紅色的組合語言會在堆疊中放入這一個部分

## 7.a x86指令，cld的意義

Clears the DF flag in the EFLAGS register. When the DF flag is set to 0, string operations increment the index registers (ESI and/or EDI).

## 7.b 自行解釋「do\_divide\_error」的部分

其實我在後面有寫...自己再稍加整理就是答案了



# 繳交方式

- 請照著這份投影片，將所有的題目做過一次，每個考題都要「截圖」或者用「手機拍照」
- 繳交期限11/21（星期三），晚上11:59:59
- 繳交方式：助教將在星期六之前公布





# 補充教材： 將「除以零」轉成 signal的方式

## 7.b 通知，在task\_struct內 設定「發生了signal」

divide\_error → do\_divide\_error → do\_error\_trap

do\_error\_trap

fill\_trap\_info

do\_trap

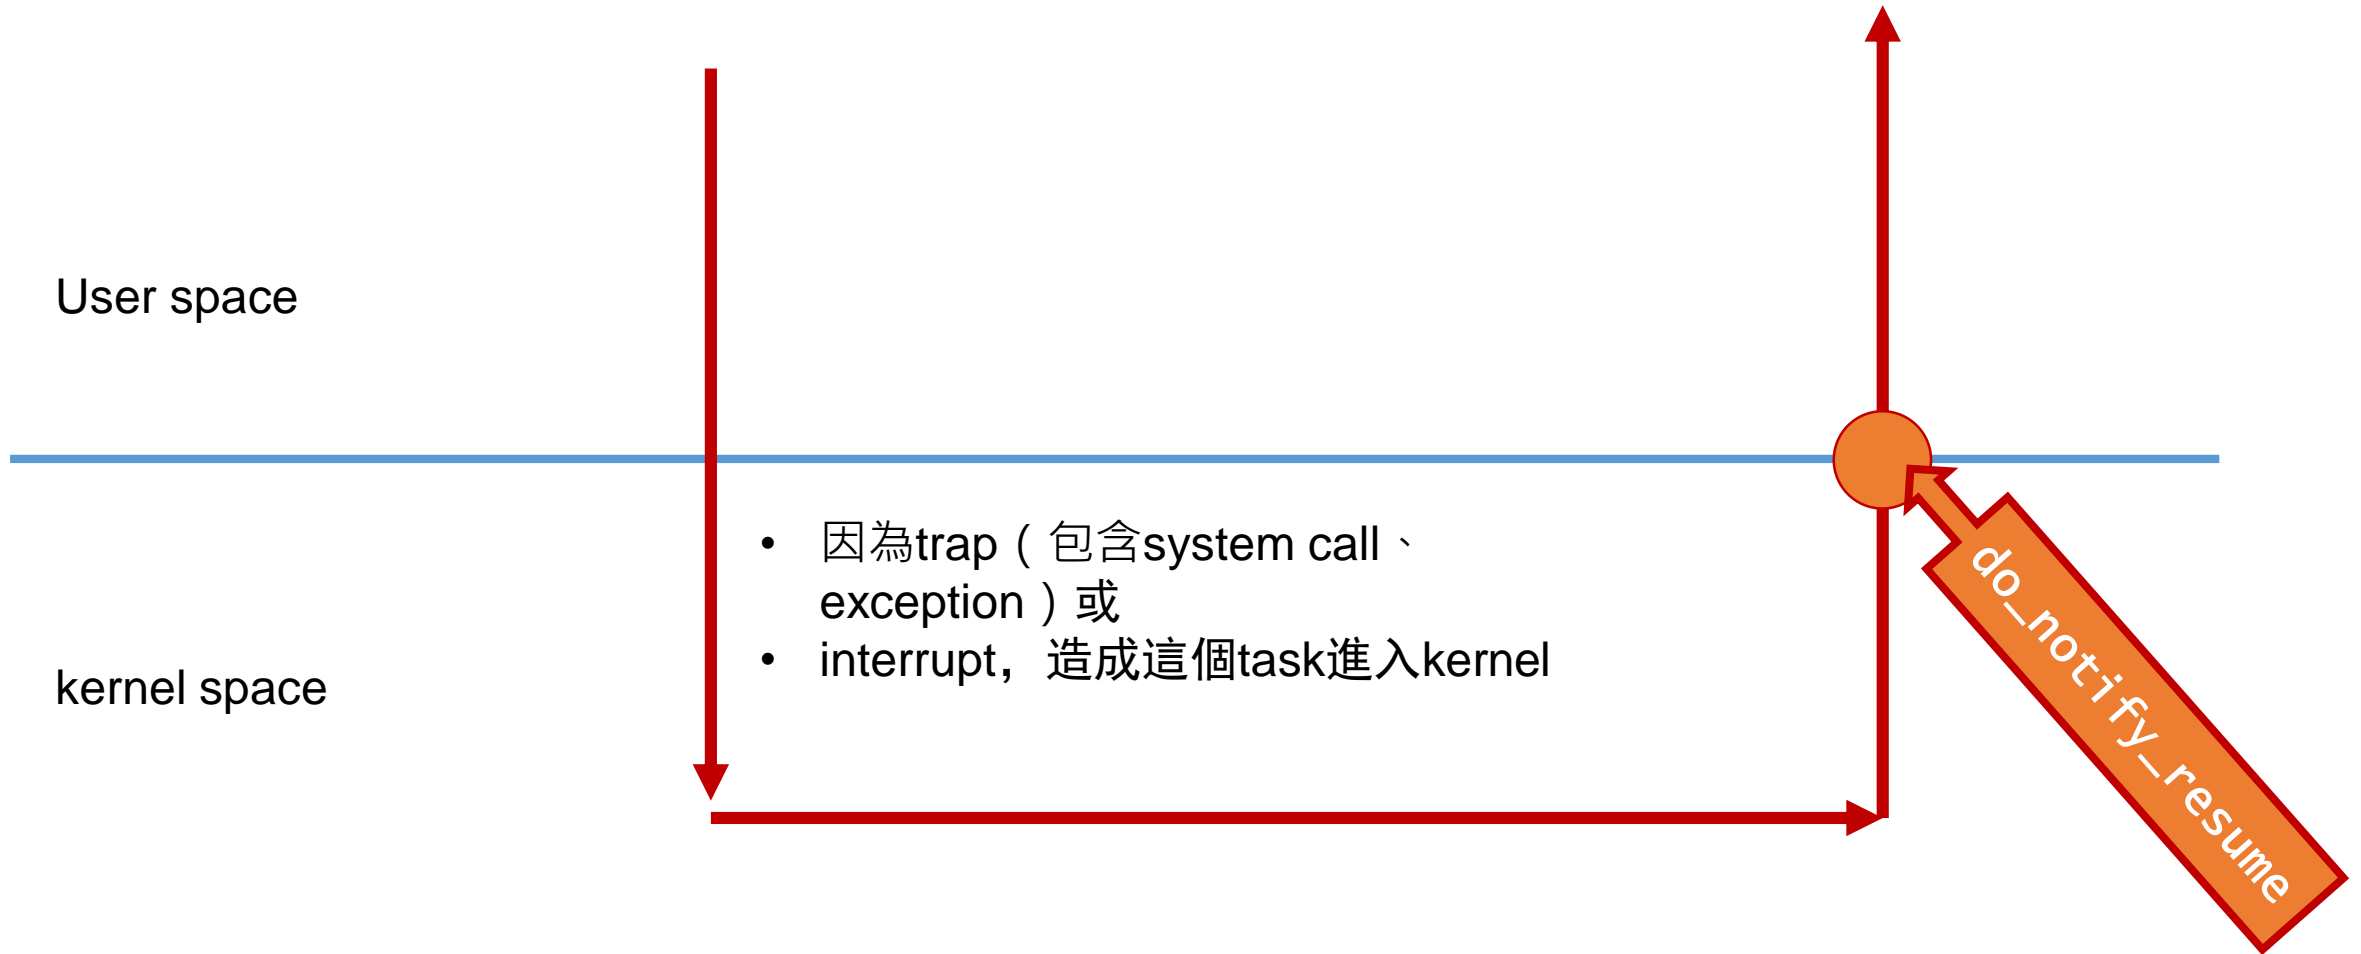
unhandled\_signal

force\_sig\_info

specific\_send\_sig\_info

\_\_send\_signal ( 在這個地方會設定flag )

## 7.b 示意圖



# do\_notify\_resmue的呼叫點

- ENTRY(native\_iret)
  - 從中斷返回時呼叫do\_notify\_resmue
- ENTRY(system\_call)
  - 從system call返回時呼叫do\_notify\_resmue
- error\_exit
  - Divide by zero，最後就是「跳到」error\_exit

# ENTRY(native\_iret)

## /arch/x86/kernel/entry\_64.S

```
1. ENTRY(native_iret)
2. /*...*/
3. retint_signal:
4. testl $_TIF_DO_NOTIFY_MASK,%edx
5. jz     retint_swapgs
6. TRACE_IRQS_ON
7. ENABLE_INTERRUPTS(CLBR_NONE)
8. SAVE_REST
9. movq $-1,ORIG_RAX(%rsp)
10.xorl %esi,%esi      # oldset
11.movq %rsp,%rdi      # &pt_regs
12.call do_notify_resume
13.RESTORE_REST
14.DISABLE_INTERRUPTS(CLBR_NONE)
15.TRACE_IRQS_OFF
16.GET_THREAD_INFO(%rcx)
17.jmp retint_with_reschedule
18./*...*/
```

## 7.b ENTRY(error\_exit)

error\_exit → retint\_careful

```
1.  retint_careful:
2.  CFI_RESTORE_STATE
3.  bt    $TIF_NEED_RESCHED,%edx
4.  jnc   retint_signal
5.  TRACE_IRQS_ON
6.  ENABLE_INTERRUPTS(CLBR_NONE)
7.  pushq_cfi %rdi
8.  SCHEDULE_USER
9.  popq_cfi %rdi
10. GET_THREAD_INFO(%rcx)
11. DISABLE_INTERRUPTS(CLBR_NONE)
12. TRACE_IRQS_OFF
13. jmp  retint_check

14. retint_signal:
15. testl $_TIF_DO_NOTIFY_MASK,%edx
16. jz    retint_swapgs
17. TRACE_IRQS_ON
18. ENABLE_INTERRUPTS(CLBR_NONE)
19. SAVE_REST
20. movq $-1,ORIG_RAX(%rsp)
21. xorl %esi,%esi      # oldset
22. movq %rsp,%rdi      # &pt_regs
23. call do_notify_resume
24. RESTORE_REST
25. DISABLE_INTERRUPTS(CLBR_NONE)
26. TRACE_IRQS_OFF
27. GET_THREAD_INFO(%rcx)
28. jmp  retint_with_reschedule
```

# ENTRY(system\_call)

## /arch/x86/kernel/entry\_64.S

```
1. ENTRY(system_call)
2. /* Either reschedule or signal or syscall exit tracking needed. */
3. /* First do a reschedule test. */
4. /* edx:    work, edi: workmask */
5. /*...*/
6. int_signal:
7. testl $_TIF_DO_NOTIFY_MASK,%edx
8. jz 1f
9. movq %rsp,%rdi          # &ptregs -> arg1
10. xorl %esi,%esi         # oldset -> arg2
11. call do_notify_resume
12. 1:    movl $_TIF_WORK_MASK,%edi
13. /*...*/
```

```
/* /arch/x86/include/asm/thread_
info.h */
/* Only used for 64 bit */
#define _TIF_DO_NOTIFY_MASK \
(_TIF_SIGPENDING | \
 _TIF_NOTIFY_RESUME | \
 _TIF_USER_RETURN_NOTIFY | \
 _TIF_UPROBE)
```

# 分析「do\_notify\_resume」

## do\_notify\_resume→do\_signal→get\_signal

```
(gdb) p ksig->info.si_signo
```

```
$1 = 8 # 8) SIGFPE
```

```
(gdb) bt
```

```
#0 get_signal (ksig=0xffff88000e5afe50) at kernel/signal.c:2358
```

```
#1 0xffffffff81005349 in do_signal (regs=0xffff88000e5aff58) at arch/x86/kernel/signal.c:703
```

```
#2 0xffffffff81005984 in do_notify_resume (regs=<optimized out>, unused=<optimized out>, thread_info_flags=<optimized out>) at arch/x86/kernel/signal.c:748
```

```
#3 <signal handler called>
```



# 附錄：關於x86硬體簡介



# 認識x86

- x86處理器，起始於1978年，8086的定址空間只有16bit（65K），大部分的電腦只有4~16K的記憶體
- 程式主要使用組合語言，例如：著名的DOS（Microsoft的作業系統）
- 為了方便程式設計師撰寫程式，組合語言的功能非常的強大，例如：有專門給字串處理用的組語
- 在記憶體方面也給程式設計師很大的方便，例如：
  - cs:0x800 ;代表程式區段第800道指令
  - ds:0x700 ;代表資料區段第800個位置的地方
  - ss:0x600 ;代表堆疊裡面第600號的位置



# 認識x86

- 下面位置中，cs、ds、ss稱之為區段，segment
  - cs:0x800 ;代表程式區段第800道指令
  - ds:0x700 ;代表資料區段第800個位置的地方
  - ss:0x600 ;代表堆疊裡面第600號的位置



# segment

- segment是一個古老的設計，但x86-64傳承了自8086起的這個古老設計
- 在x86-64原則上將segment剔除，但是還是會看到
  - Segment selector ( 一個數字，用來選擇目前要用哪一個Segment descriptor )
  - Segment descriptor ( 一個資料結構，主要記載segment的開始位置、大小、存取的屬性 )

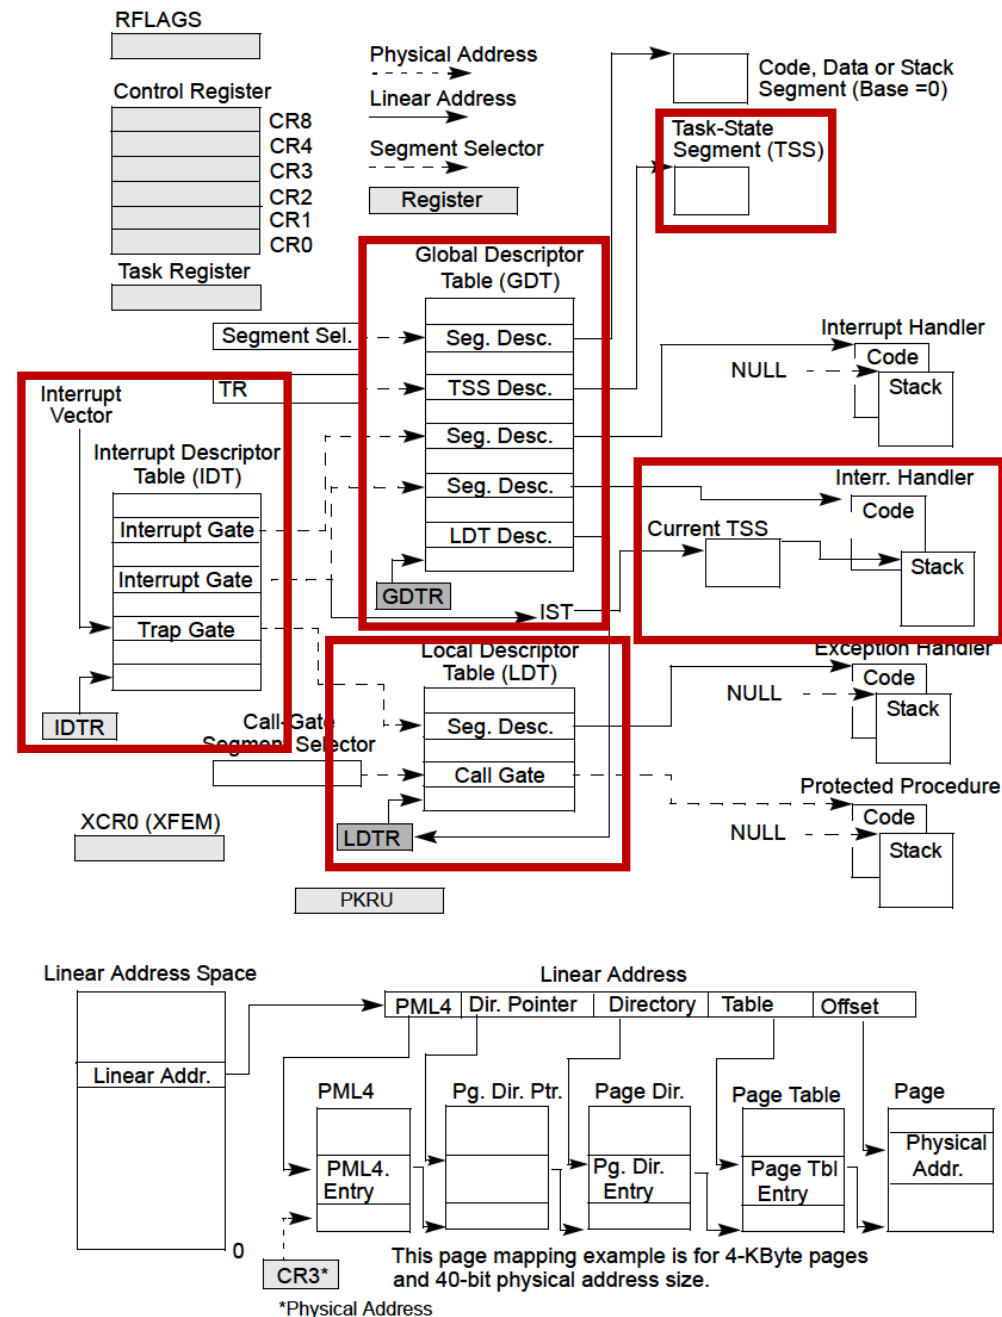


# x86-64中斷向量表



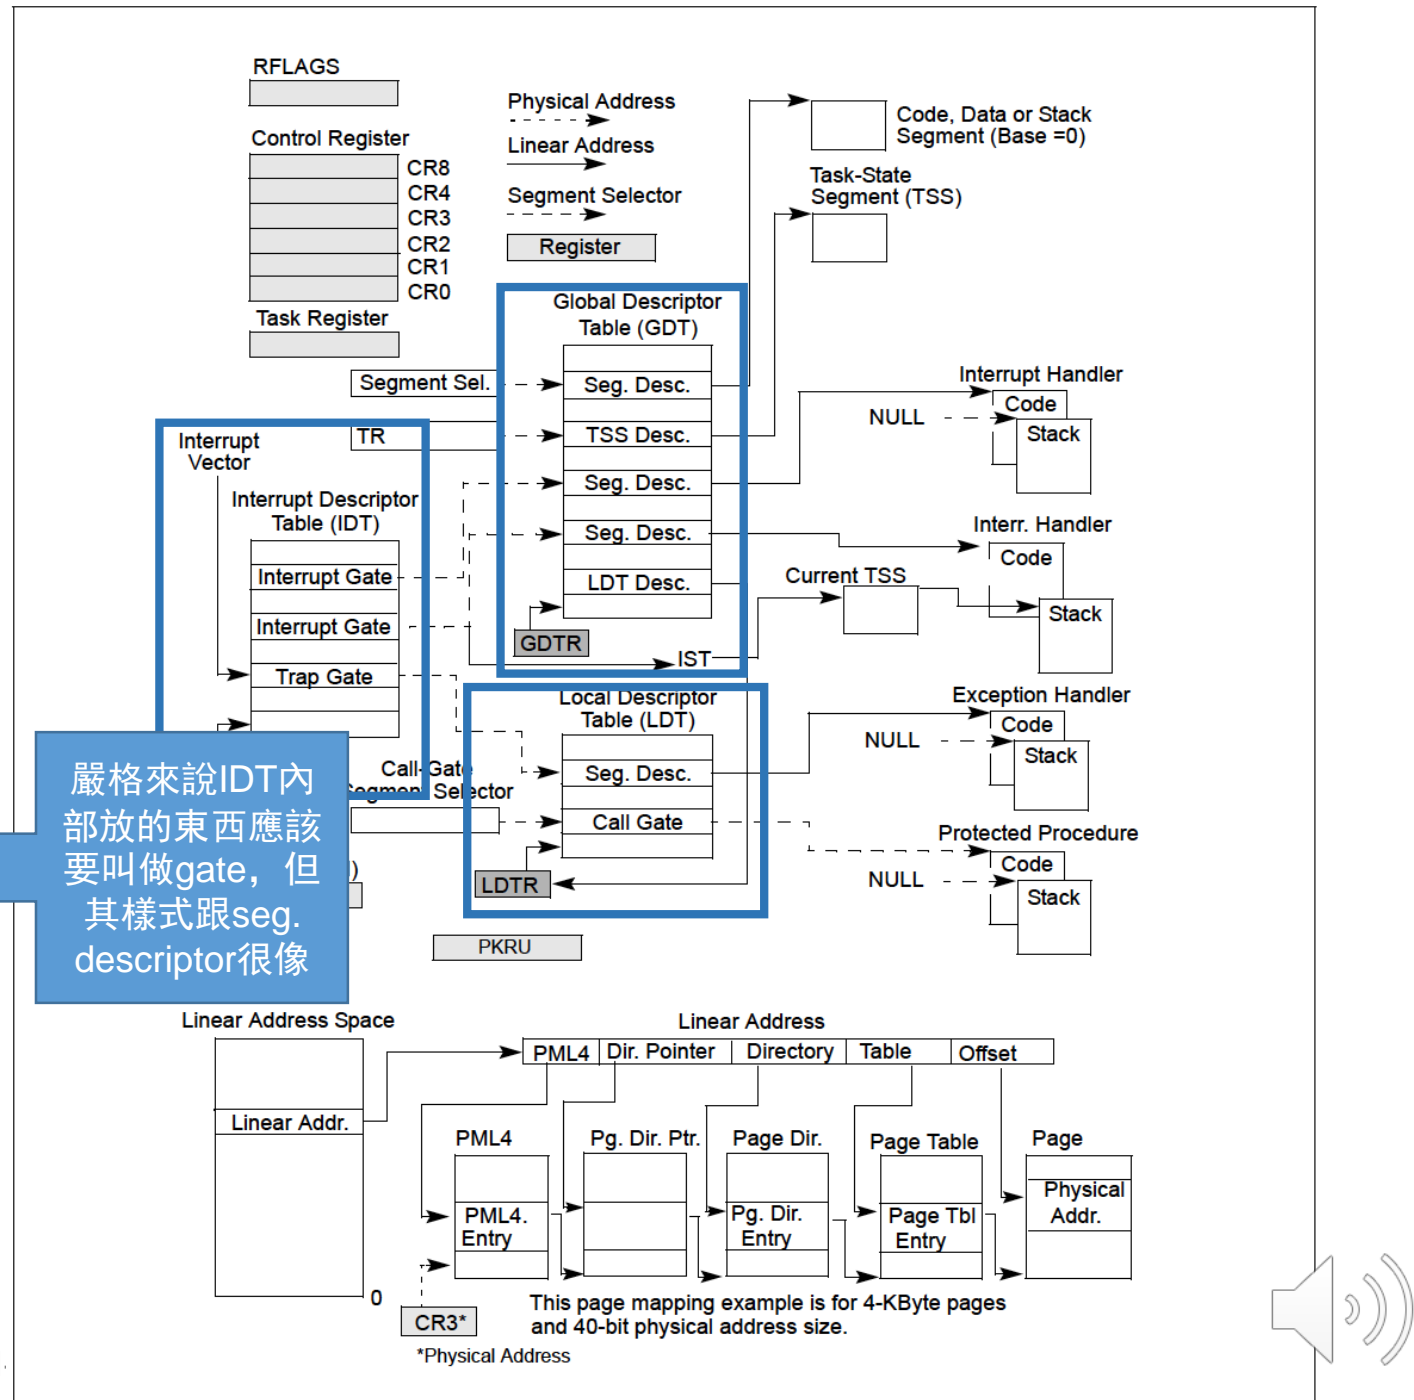
# IA-32e system level registers

紅色部分是我們會用到的部分



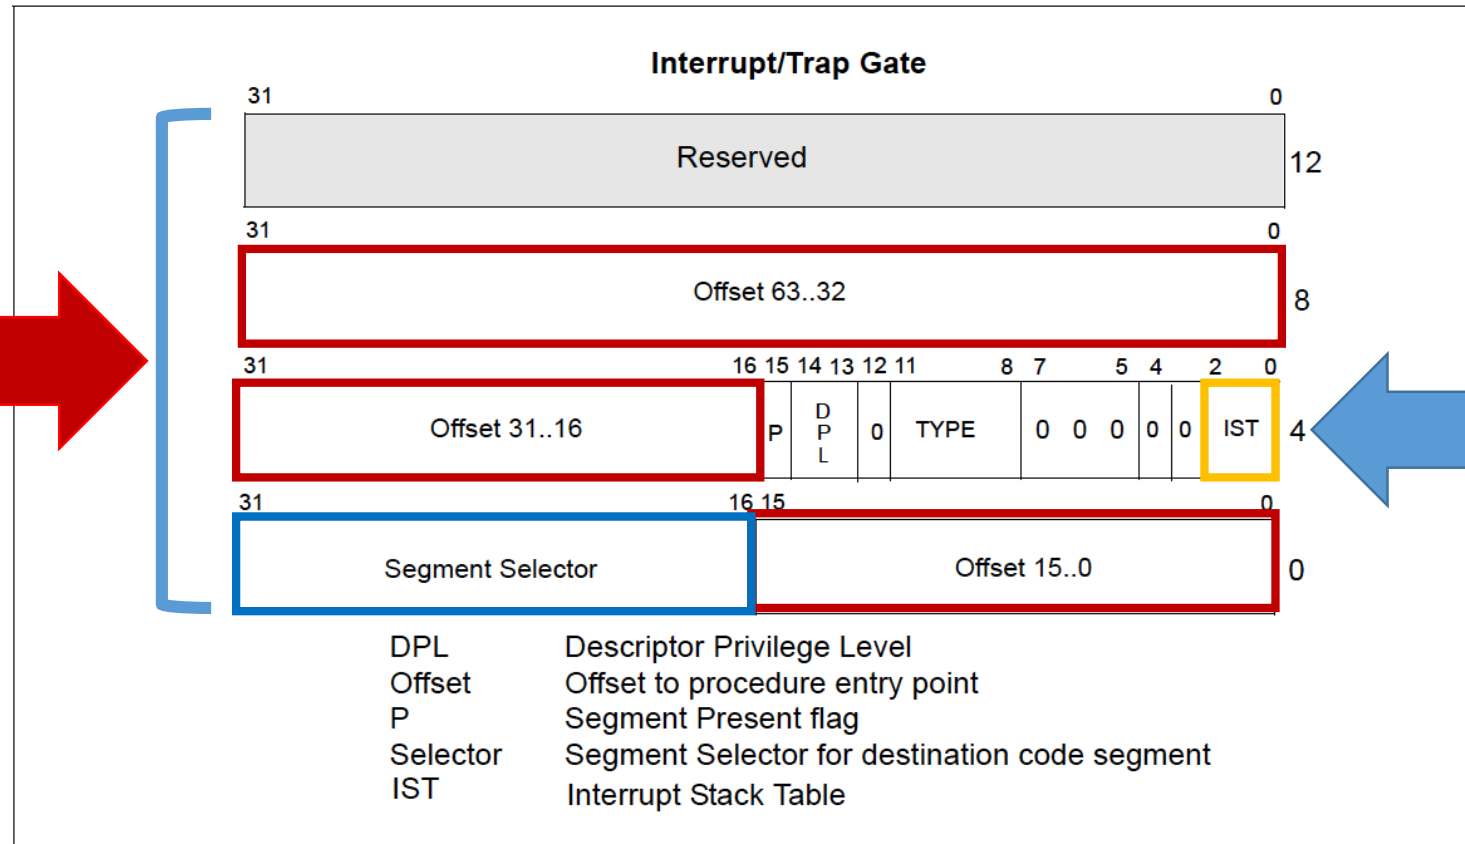
# IA-32e system level registers

- 藍色部分就是segment descriptor，系統有很多segment descriptor，以陣列的方式存在
- 在x86-64中有三個segment descriptor「陣列」，分別用IDTR、GDTR、LDTR指向這三個陣列
- segment selector就是index用來指向LDTR、GDTR，而IDTR陣列，就是中斷向量表



# IDT Gate Descriptor

中斷向量表的其中一個entry，大小為128 bits，紅色的部分共64 bits指向ISR



IST，值介於1~7，可以讓OS選擇發生中斷的時候要切換到哪一個stack

ISR : interrupt server routine



Figure 6-7. 64-Bit IDT Gate Descriptors



# Hardware-generated exceptions

INT_NUM	Short Description PM
<b>0x00</b>	<a href="#">Division by zero</a>
0x01	Single-step interrupt (see <a href="#">trap flag</a> )
0x02	<a href="#">NMI</a>
0x03	Breakpoint (callable by the special 1-byte instruction 0xCC, used by debuggers)
0x04	Overflow
0x05	Bounds
0x06	Invalid Opcode
0x07	Coprocessor not available
0x08	<a href="#">Double fault</a>
0x09	Coprocessor Segment Overrun ( <i>386 or earlier only</i> )
0x0A	Invalid Task State Segment
0x0B	Segment not present
0x0C	Stack Fault
0x0D	<a href="#">General protection fault</a>
0x0E	<a href="#">Page fault</a>
0x0F	<i>reserved</i>
0x10	Math Fault
0x11	Alignment Check
0x12	Machine Check
0x13	<a href="#">SIMD</a> Floating-Point Exception
0x14	Virtualization Exception
0x15	Control Protection Exception



# /arch/x86/include/asm/traps.h

```
1.  enum {
2.      X86_TRAP_DE = 0,      /* 0, Divide-by-zero */
3.      X86_TRAP_DB,          /* 1, Debug */
4.      X86_TRAP_NMI,         /* 2, Non-maskable Interrupt */
5.      X86_TRAP_BP,          /* 3, Breakpoint */
6.      X86_TRAP_OF,          /* 4, Overflow */
7.      X86_TRAP_BR,          /* 5, Bound Range Exceeded */
8.      X86_TRAP_UD,          /* 6, Invalid Opcode */
9.      X86_TRAP_NM,          /* 7, Device Not Available */
10.     X86_TRAP_DF,           /* 8, Double Fault */
11.     X86_TRAP_OLD_MF,       /* 9, Coprocessor Segment Overrun */
12.     X86_TRAP_TS,           /* 10, Invalid TSS */
13.     X86_TRAP_NP,           /* 11, Segment Not Present */
14.     X86_TRAP_SS,           /* 12, Stack Segment Fault */
15.     X86_TRAP_GP,           /* 13, General Protection Fault */
16.     X86_TRAP_PF,           /* 14, Page Fault */
17.     X86_TRAP_SPURIOUS,     /* 15, Spurious Interrupt */
18.     X86_TRAP_MF,           /* 16, x87 Floating-Point Exception */
19.     X86_TRAP_AC,           /* 17, Alignment Check */
20.     X86_TRAP_MC,           /* 18, Machine Check */
21.     X86_TRAP_XF,           /* 19, SIMD Floating-Point Exception */
22.     X86_TRAP_IRET = 32,    /* 32, IRET Exception */
23. };
```



# 對IDT進一步解釋 先解釋segment

# gate\_struct64

## /arch/x86/include/asm/desc\_defs.h

```
1.  /* 16byte gate */
2.  struct gate_struct64 {
3.      u16 offset_low;
4.      u16 segment;
5.      unsigned ist : 3, zero0 : 5, type : 5, dpl : 2, p : 1;
6.      u16 offset_middle;
7.      u32 offset_high;
8.      u32 zero1;
9.  } __attribute__((packed));
```

# segments in x64

In 64-bit mode: CS, DS, ES, SS are treated as if each **segment base is 0**, regardless of the value of the associated segment descriptor base. This creates a **flat address space for code, data, and stack**. FS and GS are exceptions. Both segment registers may be used as additional base registers in linear address calculations (in the addressing of local data and certain operating system data structures).

# segments in x64

- 大致上意思是說：
  - segment原則上沒有用處
  - GS和FS可能會用到

# segment descriptor

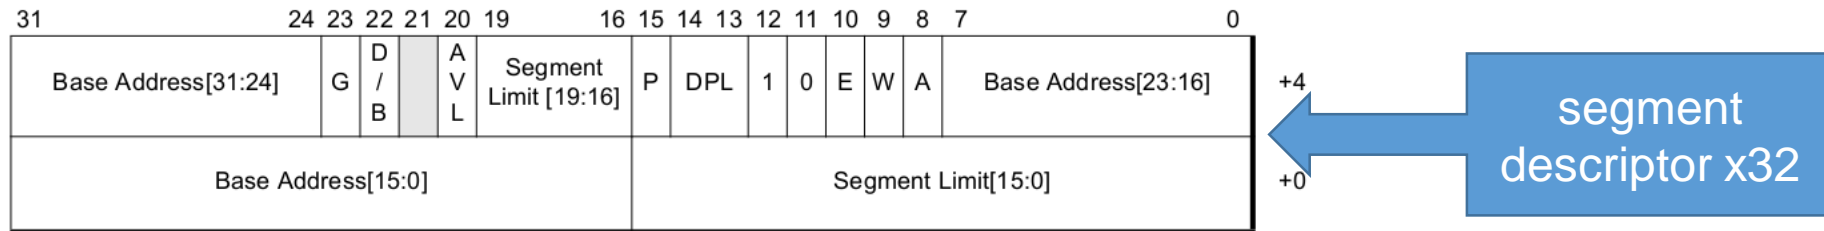


Figure 4-15. Data-Segment Descriptor—Legacy Mode

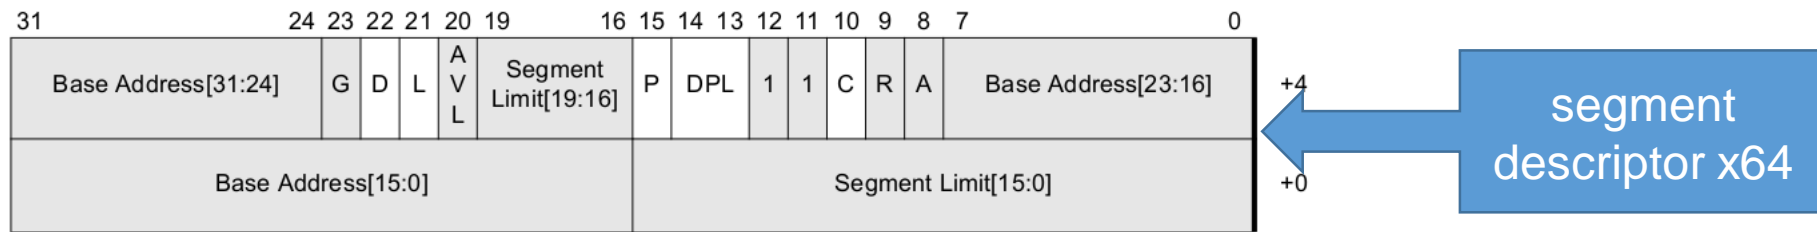



Figure 4-20. Code-Segment Descriptor—Long Mode

# desc\_struct

## /arch/x86/include/asm/desc\_defs.h

```
1. struct desc_struct {  
2.     union {  
3.         struct {  
4.             unsigned int a;  
5.             unsigned int b;  
6.         };  
7.         struct {  
8.             u16 limit0;  
9.             u16 base0;  
10.            unsigned base1: 8, type: 4, s: 1, dp1: 2, p: 1;  
11.            unsigned limit: 4, avl: 1, l: 1, d: 1, g: 1, base2: 8;  
12.        };  
13.    };  
14. } __attribute__((packed));
```





# 對IDT進一步解釋 再解釋IST

# Task state segment ( tss )

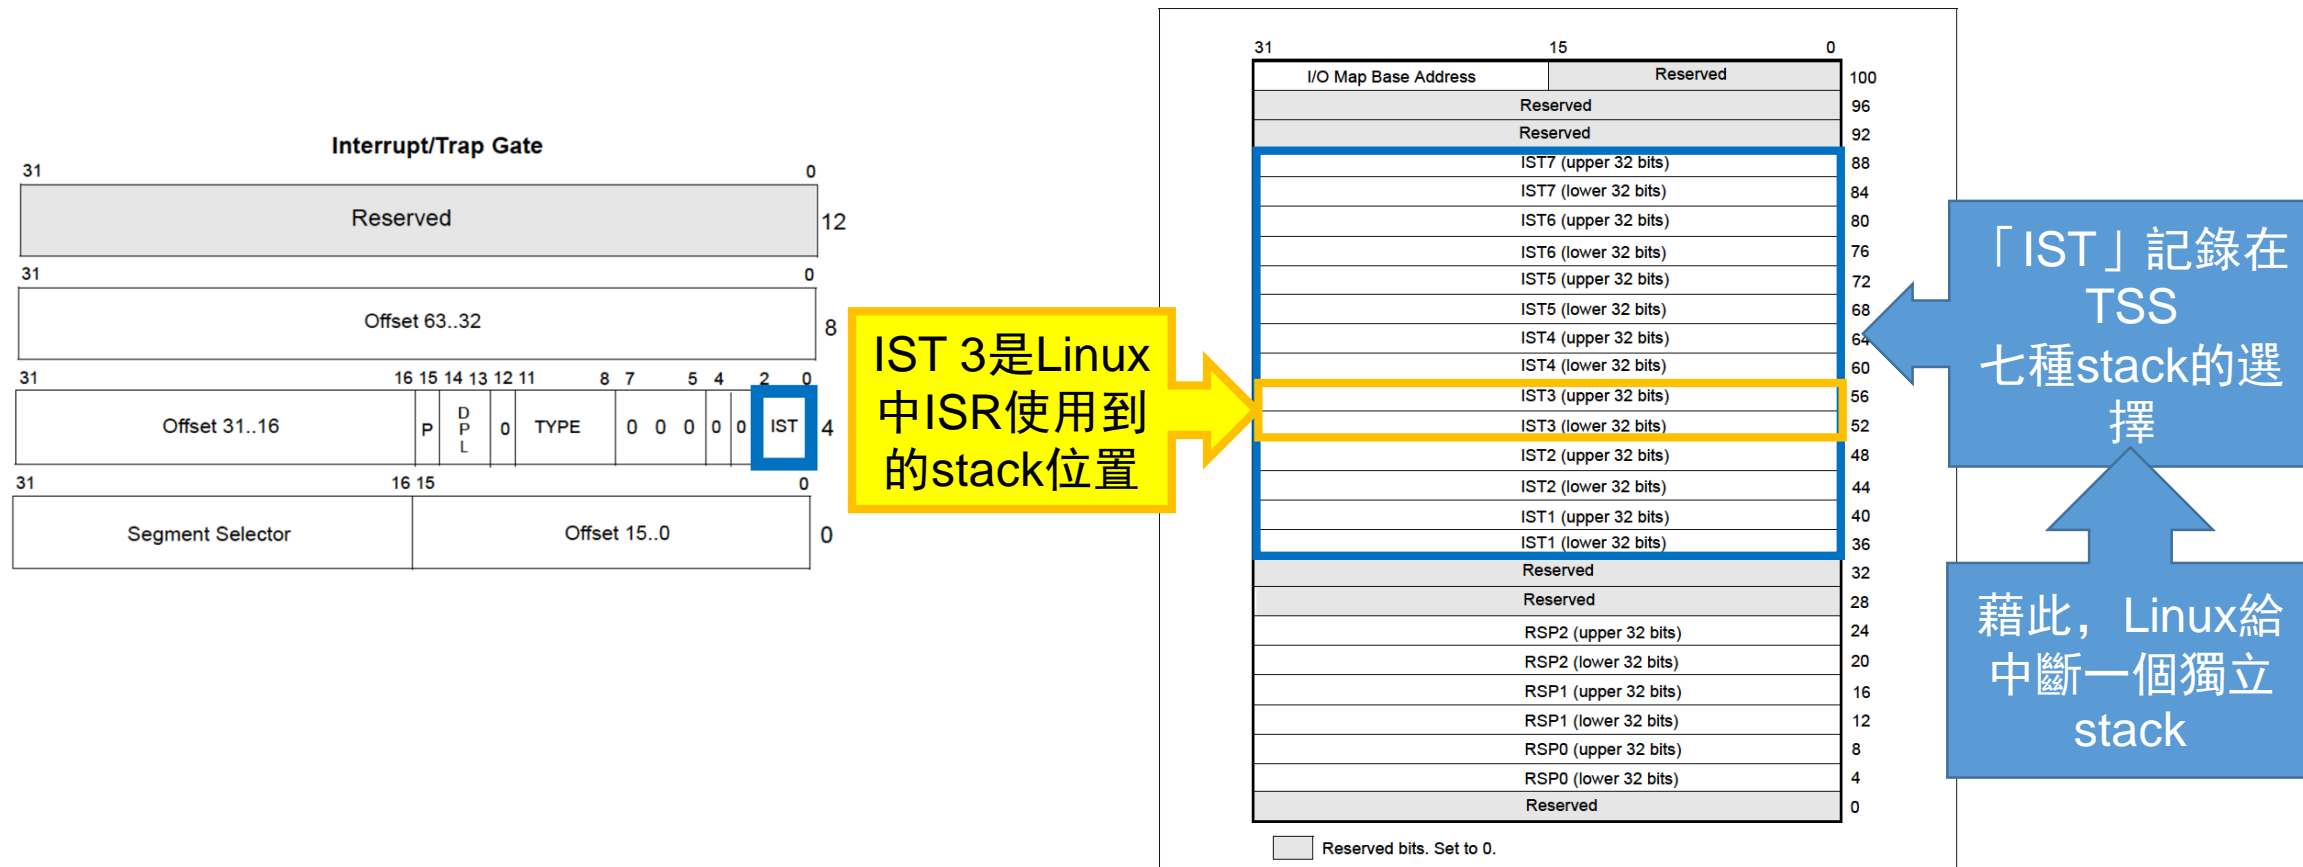


Figure 7-11. 64-Bit TSS Format



# tss\_struct

## /arch/x86/include/asm/processor.h

```
1. struct tss_struct {
2.     /*
3.      * The hardware state:
4.      */
5.     struct x86_hw_tss    x86_tss;
6.     unsigned long        io_bitmap[IO_BITMAP_LONGS + 1];
7.     unsigned long        stack[64];
8. } ____cacheline_aligned;
```



# tss : initialization

## /arch/x86/include/asm/processor.h

```
1. /*
2. * per-CPU TSS segments. Threads are completely 'soft' on Linux,
3. * no more per-task TSS's. The TSS size is kept cacheline-aligned
4. * so they are allowed to end up in the .data..cacheline_aligned
5. * section. Since TSS's are completely CPU-local, we want them
6. * on exact cacheline boundaries, to eliminate cacheline ping-pong.
7. */
8. __visible DEFINE_PER_CPU_SHARED_ALIGNED(struct tss_struct, init_tss) = INIT_TSS;
```



# tss : initialization

/arch/x86/include/asm/processor.h

```
1. #define INIT_TSS {
2.     .x86_tss = {
3.         .sp0      = sizeof(init_stack) + (long)&init_stack, \
4.         .ss0      = __KERNEL_DS, \
5.         .ss1      = __KERNEL_CS, \
6.         .io_bitmap_base = INVALID_IO_BITMAP_OFFSET, \
7.     }, \
8.     .io_bitmap    = { [0 ... IO_BITMAP_LONGS] = ~0 }, \
9. }
```



# tss in ctw\_sw

## /arch/x86/kernel/process\_64.c

```
1. __visible __notrace_funcgraph struct task_struct *
2. __switch_to(struct task_struct *prev_p, struct task_struct *next_p)
3. {
4.     int cpu = smp_processor_id();
5.     struct tss_struct *tss = &per_cpu(init_tss, cpu);
6.     /*...*/
7.
8.     if (unlikely(task_thread_info(next_p)->flags & _TIF_WORK_CTXSW_NEXT ||
9.             task_thread_info(prev_p)->flags & _TIF_WORK_CTXSW_PREV))
10.         __switch_to_extra(prev_p, next_p, tss);
11. }
```



# tss in ctw\_sw

## /arch/x86/kernel/process.c

```
1. void __switch_to_xtra(struct task_struct *prev_p, struct task_struct *next_p,
2.                       struct tss_struct *tss)
3. /*...*/
4. if (test_tsk_thread_flag(next_p, TIF_IO_BITMAP)) {
5.     memcpy(tss->io_bitmap, next->io_bitmap_ptr,
6.           max(prev->io_bitmap_max, next->io_bitmap_max));
7. } else if (test_tsk_thread_flag(prev_p, TIF_IO_BITMAP)) {
8.     memset(tss->io_bitmap, 0xff, prev->io_bitmap_max);
9. }
10. /*...*/
11. }
```



# trap\_init

- call `cpu_init()` to `do`:
- initialize per-CPU state
- reload the GDT and IDT
- mask off the eflags `NT` (Nested Task) bit
- set up and load the per-CPU TSS and LDT
- clear 6 debug registers (0, 1, 2, 3, 6, and 7)
- `stts()`: set the `0x08` bit (TS: Task Switched) in CR0 to enable lazy
- `register` saves on context switches





# trap\_init → cpu\_init

## /arch/x86/kernel/cpu/common.c

```
1. void cpu_init(void) {
2.     me = current;
3.     syscall_init();
4.     barrier();
5.     //set up and load the per-CPU TSS
6.     if (!oist->ist[0]) {
7.         char *estacks = per_cpu(exception_stacks, cpu);
8.         for (v = 0; v < N_EXCEPTION_STACKS; v++) {
9.             estacks += exception_stack_sizes[v];
10.            oist->ist[v] = t->x86_tss.ist[v] =
11.                (unsigned long)estacks;
12.            if (v == DEBUG_STACK-1)
13.                per_cpu(debug_stack_addr, cpu) = (unsigned long)estacks;
14.        }
15.    }
16.    t->x86_tss.io_bitmap_base = offsetof(struct tss_struct, io_bitmap);
}
```



# syscall & sysret

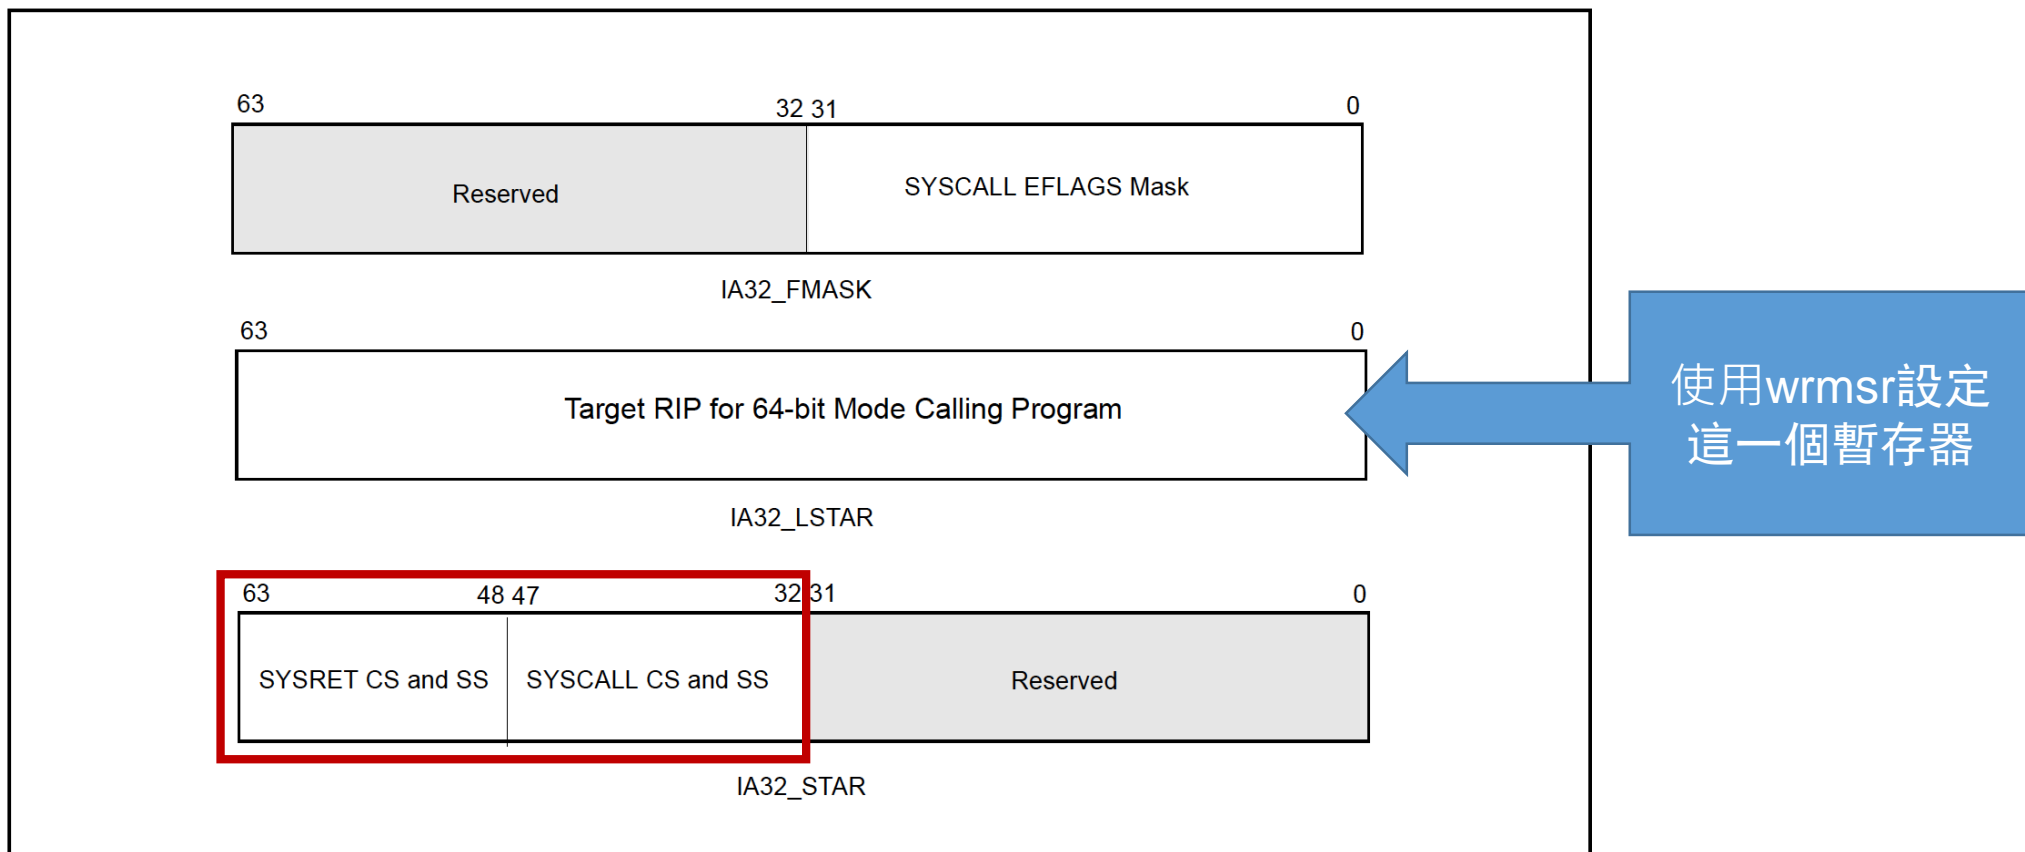


Figure 5-14. MSRs Used by SYSCALL and SYSRET



trap\_init → cpu\_init → syscall\_init  
[/arch/x86/kernel/cpu/common.c](#)

```
1. void syscall_init(void)
2. {
3.     wrmsrl(MSR_STAR, ((u64)__USER32_CS)<<48 | ((u64)__KERNEL_CS)<<32);
4.     wrmsrl(MSR_LSTAR, system_call);
5.     wrmsrl(MSR_CSTAR, ignore_sysret);
6.     /* Flags to clear on syscall */
7.     wrmsrl(MSR_SYSCALL_MASK,
8.             X86_EFLAGS_TF|X86_EFLAGS_DF|X86_EFLAGS_IF|
9.             X86_EFLAGS_IOPL|X86_EFLAGS_AC|X86_EFLAGS_NT);
10. }
```

